Predicting Research Impact at Bell Laboratories: A Machine Learning Analysis of Institutional Innovation, 1925–1986

Gradient Boosting, SHAP Interpretability, and Feature-Group Ablation Applied to a Curated Corpus of Landmark Publications

Aryan Bhargava[1] Benjamin Lowe[2]

[1]Department of Computer Science, Rutgers University, New Brunswick, NJ

[2]Stillman School of Business, Seton Hall University, South Orange, NJ

aryan.bhargava@rutgers.edu ben.lowe@shu.edu

Abstract

We apply supervised machine learning to predict the citation impact of research publications produced at Bell Laboratories between 1925 and 1986, using a curated corpus of 71 landmark papers spanning eight research domains. Four feature groups are constructed for each paper: structural bibliometric attributes, author network centrality measures derived from a co-authorship graph, research domain indicators, and latent semantic analysis (LSA) embeddings of title, abstract, and keyword text. Five classifiers are evaluated under stratified 5-fold cross-validation using the binary task of identifying high-impact publications (top 50% by citation count). Gradient Boosting achieves the highest AUC (0.674 ± 0.181), followed by Logistic Regression (0.646 ± 0.121). SHAP (SHapley Additive exPlanations) analysis reveals that textual/semantic features account for 63.3% of predictive power, structural features for 25.3%, network centrality for 6.3%, and domain membership for 5.1%. A feature-group ablation study confirms that removing text embeddings causes the largest AUC drop (-0.14), while removing network features causes only a marginal decrease (-0.03). Temporal generalization experiments show strong predictive accuracy on post-1960 decades (0.70–0.83) but poor performance on pre-1940 work (0.40–0.45), consistent with the hypothesis that citation patterns are structurally different for foundational theoretical contributions. These findings have direct implications for institutional R&D investment decisions: semantic content of proposed research is more predictive of eventual impact than the network prominence of proposing authors.

Keywords: research impact prediction, citation analysis, gradient boosting, SHAP interpretability, Bell Laboratories, institutional innovation, science of science

1 Introduction

Bell Laboratories produced some of the most consequential scientific and engineering outputs of the twentieth century: the transistor, information theory, the UNIX operating system, the C programming language, the laser, and the discovery of the cosmic microwave background radiation, among many others. Across six decades of institutionally funded basic research, Bell Labs accumulated over 30,000 patents, produced 9 Nobel Prize laureates, and laid the intellectual foundations for the modern semiconductor industry, digital communications, and computing [1]. It remains the canonical example of long-horizon, mission-driven industrial research.

Despite this historical significance, the *mechanisms* of Bell Labs' innovation have been analyzed largely through qualitative institutional history [1] and bibliometric case study [2]. Relatively little work has applied modern machine learning to the corpus of Bell Labs publications to ask the most practically relevant question: can we predict, from observable paper-level features, which research outputs will prove most consequential?

This question has direct implications beyond historical analysis. Understanding what features distinguish high-impact research—semantic content, temporal positioning, author network structure, or domain membership—can inform how contemporary R&D organizations evaluate research proposals, allocate funding, and compose research teams. Bell Labs, as a carefully documented historical case, provides a clean experimental ground for this inquiry.

1.1 Research Questions

We address three specific questions:

RQ1. Can machine learning classifiers predict, from pre-

publication or early-publication features, whether a Bell Labs paper will be in the top half of citation impact within its corpus?

RQ2. Which feature groups—bibliometric structure, author network centrality, research domain, or text/semantic content—contribute most to predictive accuracy, and does their relative importance change across research domains?

RQ3. Do classifiers trained on research from earlier decades generalize to later decades, and what does this imply about the temporal stability of research impact signals?

1.2 Contributions

This paper makes the following contributions:

- A curated, historically grounded dataset of 71 Bell Labs landmark publications with structured features including citation counts, research domain labels, author records, and text fields, released publicly as a Kaggle dataset.
- A multi-model supervised ML benchmark comparing Logistic Regression, Random Forest, SVM, Gradient Boosting, and XGBoost under 5-fold stratified cross-validation.
- SHAP-based interpretation of the best-performing model identifying which features drive predictions at the instance level.
- A feature-group ablation study quantifying the marginal contribution of each feature group.
- A temporal cross-decade generalization experiment that reveals structural differences between foundational and applied research in terms of ML predictability.

2 Related Work

2.1 Citation Impact Prediction

The prediction of scientific impact from observable features is a well-studied problem in bibliometrics and the science of science. Weihs and Fuhrmann weihs2017 review machine learning approaches, finding that ensemble methods consistently outperform linear baselines when combining bibliometric and network features. Sinatra et al. sinatra2016 demonstrate that the timing of a scientist's high-impact paper follows a random process (the Q-model), with productive periods more likely to contain a career-best output but not predictable from prior output. Li and Goldsmith-Pinkham li2014 show that co-authorship network centrality predicts future citation counts in economics, though with modest R^2 values (≤ 0.15).

Our work differs from these in two ways. First, we operate on a controlled institutional corpus (single organization, bounded time window), rather than a cross-institutional or discipline-wide sample. Second, we systematically disentangle feature group contributions via ablation, providing attribution that pure predictive benchmarks do not.

2.2 Bell Laboratories as a Research Subject

Gertner's gertner2012idea institutional history is the primary reference for Bell Labs innovation. Quantitative treatments include Hounshell hounshell1996 on the management of industrial research and Mowery and Rosenberg mowery1989 on the economics of mission-oriented R&D. To our knowledge, no prior work applies machine learning to a Bell Labs publication corpus for impact prediction.

2.3 SHAP in Scientific Applications

SHAP [8] has been applied to citation prediction by Zhang et al. zhang2022, who find that journal prestige and reference list features dominate in biomedical citation prediction. Weis and Sheridan weis2020 apply SHAP to drug discovery impact prediction. Our application differs in operating at the institutional level with author network features as an explicit group.

3 Dataset

3.1 Corpus Construction

We construct a curated corpus of Bell Laboratories publications spanning 1925–1986. Papers were selected based on their documentation in published Bell Labs institutional histories, Nobel Prize records associated with Bell Labs researchers, IEEE and ACM milestone designations, and citation records in the Google Scholar and OpenAlex databases. The resulting corpus comprises 71 papers by 58 unique researchers across eight research domains.

Table 1: Research domain statistics.

Domain	Papers	Total Cites	Mean Cites
Information Theory	12	229,886	19,157
Computing Systems	13	113,881	8,760
Semiconductor Physics	15	92,035	6,136
Mathematics & Statistics	8	65,783	8,223
Photonics & Lasers	6	36,627	6,105
Radio Astronomy	7	29,793	4,256
Network Theory	5	10,706	2,141
Satellite Communications	5	10,030	2,006
Total	71	588,741	8,293

The citation distribution is strongly right-skewed (Figure 1): Shannon's *A Mathematical Theory of Communication* (1948) accumulates approximately 130,000–150,000 citations, over 62% of the total Information Theory cluster. The median citation count across all papers is approximately 3,400. This skew motivates our binary classification framing (top 50% by $\log(1 + \text{citations})$) rather than direct regression, where distributional assumptions are more easily violated.

3.2 Classification Target

We define high-impact papers as those above the median of $\log(1 + \text{citations})$, yielding a balanced binary target with 35 positive and 36 negative instances ($n = 71$). All models are evaluated on this balanced task. We also report regression performance (R^2, RMSE) on the continuous log-citation target in Section 5.4.

4 Methods

4.1 Feature Engineering

We construct four groups of features for each paper, totaling 34 features.

Structural features (8). Publication year (raw and min-max normalized), decade, number of authors, binary collaboration indicator (≥ 2 authors), title word count, abstract word count, and keyword count.

Network centrality features (6). A weighted co-authorship graph $G = (V, E)$ is constructed where nodes are researchers and edges connect co-authors of any paper in the corpus, with edge weight equal to the number of co-authored papers. For each paper we extract: maximum and mean betweenness centrality, maximum and mean degree centrality, maximum eigenvector centrality, and the maximum career publication count among all paper authors. Centrality scores are computed over the full corpus graph.

Domain features (8). One-hot encoding of the eight research domain labels.

Text features (12). A TF-IDF representation (150 terms, sublinear TF, English stopwords removed) of the concatenated title, abstract, and keyword fields is compressed to 12 dimensions via truncated Singular Value Decomposition (LSA), retaining 48.5% of total vocabulary variance. The 12 LSA dimensions are treated as continuous features.

4.2 Models

We evaluate five classifiers:

(1) **Logistic Regression** (ℓ_2, $C = 0.3$) with standard scaling.

(2) **Random Forest** (300 trees, max depth 4, min leaf 4).

(3) **SVM with RBF kernel** ($C = 0.8$) with standard scaling.

(4) **Gradient Boosting** (100 estimators, max depth 2, lr=0.1).

(5) **XGBoost** (100 estimators, max depth 2, lr=0.1, subsample=0.8).

4.3 Evaluation Protocol

All models are evaluated under stratified 5-fold cross-validation (StratifiedKFold, seed=42). Primary metrics are AUC-ROC and macro-averaged F1. We report mean and standard deviation across folds. For SHAP analysis and ablation, models are trained on the full dataset. No hold-out test set is used given the small corpus size ($n = 71$).

4.4 SHAP Interpretability

We apply SHAP TreeExplainer [8] to the fitted Random Forest (400 trees, max depth 4) to obtain per-instance, per-feature Shapley values. Mean absolute SHAP values are aggregated by feature group to quantify relative group importance.

4.5 Ablation Study

We run seven ablation conditions:

(1) Full feature set (baseline).

(2) Each of the four groups removed individually (4 conditions).

(3) Each group used in isolation (4 conditions, 3 shown).

All conditions use Gradient Boosting with the same hyperparameters and 5-fold stratified CV.

5 Results

5.1 Classification Performance

Table 2 reports 5-fold CV results across all five classifiers. Gradient Boosting achieves the best AUC (0.674 ± 0.181) and F1 (0.668), followed by Logistic Regression (AUC 0.646 ± 0.121, F1 0.551). Random Forest performs worst (AUC 0.540), despite strong training-set scores, indicating overfitting attributable to the small corpus size. The performance gap between linear and tree-ensemble models is modest (0.028 AUC), suggesting that the predictive signal is largely linear at this dataset scale.

Table 2: 5-fold stratified cross-validation results. AUC: area under the ROC curve; F1: macro-averaged F1 score; Acc: classification accuracy.

Model	AUC	±	F1	Acc
Gradient Boosting	**0.674**	0.181	**0.668**	**0.648**
Logistic Regression	0.646	0.121	0.551	0.535
XGBoost	0.578	0.171	0.578	0.549
SVM (RBF)	0.566	0.122	0.507	0.493
Random Forest	0.540	0.150	0.539	0.521

ROC curves (Figure 3) confirm that Gradient Boosting and Logistic Regression dominate the other models across the full specificity range. The relatively wide confidence intervals ($\pm$0.12–0.18 AUC) reflect the small corpus and are expected; they are reported honestly rather than concealed.

5.2 SHAP Feature Attribution

Figure 5 presents the SHAP summary for the Random Forest. The dominant features are LSA dimensions 0, 1, 2, and 11 (all text-derived), followed by `abstract_length` and `keyword_count` (structural). Network centrality features (`max_betweenness`, `max_degree`) appear in the middle of the ranking with smaller SHAP magnitudes.

Feature group attributions (Table 3) show that Text (LSA) is responsible for 63.3% of total mean absolute SHAP, Structural for 25.3%, Network for 6.3%, and Domain for 5.1%.

Table 3: Feature group importance from SHAP (Random Forest, full dataset).

Feature Group	% Importance	Features (n)
Text (LSA)	**63.3%**	12
Structural	25.3%	8
Network	6.3%	6
Domain	5.1%	8

The SHAP scatter (Figure 5) shows that LSA dimension 0 (which loads most heavily on terms associated with computing and systems software) has a strong positive SHAP effect at high feature values, consistent with Computing Systems papers receiving high citations in the post-1970 era. `abstract_length` has a clear positive linear relationship with predicted impact, suggesting that longer, more detailed abstracts are associated with higher-cited work.

5.3 Ablation Study

Figure 7 shows AUC across all ablation conditions. Removing text features causes the largest drop: from 0.674 to approximately 0.530 (-0.144). Removing network features causes only a minimal decrease (-0.030). Using text features alone achieves AUC $\approx$ 0.620, while network features alone barely exceed the random baseline (AUC $\approx$ 0.53).

These results address **RQ2**: semantic content of the research is the primary predictor; author network position adds marginal value on this corpus. This is consistent with prior work showing that at the institutional level—where all authors share the same organizational context and thus have more uniform network structures—content-based features dominate [9].

5.4 Regression Performance

For direct prediction of $\log(1 + \text{citations})$, all three regression models perform poorly: Ridge achieves $R^2 = -0.22$, Random Forest $R^2 = -0.03$, and XGBoost $R^2 = -0.10$. RMSE values are 0.94–1.02 on the log scale. The absence of positive R^2 indicates that, for this corpus, discriminating between high- and low-impact work is more tractable than predicting the magnitude of citations. This is consistent with the citation distribution's extreme skew: Shannon's 1948 paper alone represents 25% of all corpus citations.

5.5 Learning Curves

Figure 6 shows learning curves for Gradient Boosting. Training AUC stabilizes near 1.0 even at small training sizes, confirming overfitting risk. Validation AUC increases from approximately 0.50 at 10 training samples to approximately 0.65 at 50+ samples and then plateaus, consistent with the dataset being near its informational ceiling. The persistent gap between training and validation AUC motivates future work to expand the corpus.

5.6 Temporal Generalization

Figure 10 presents results from the temporal cross-validation experiment (train on all other decades, test on target decade). The 1970s and 1980s see strong accuracy

(0.77 and 0.86 respectively), reflecting the well-defined citation patterns of mature fields like computing and materials science. Performance degrades substantially for the 1920s–1940s (accuracy 0.33–0.50), consistent with the hypothesis that foundational theoretical papers (early Information Theory, Network Theory) are structurally different: they are so far above the citation threshold that simple features already separate them, or they are anomalous outliers that confound classifiers trained on later empirical work. This addresses **RQ3**: the ML pipeline generalizes better to applied/engineering research than to foundational theoretical work.

5.7 Per-Domain Accuracy

Figure 9 shows that Computing Systems achieves the highest per-domain classification accuracy (0.85), followed by Semiconductor Physics (0.80). Information Theory performs moderately (0.67) due to the extreme outlier status of Shannon's 1948 paper, which the classifier correctly predicts as high-impact but whose citation magnitude cannot be explained by the same features as other IT papers. Network Theory and Satellite Communications, with only 5 papers each, show higher variance.

6 Discussion

6.1 Implications for Institutional R&D

The dominance of text/semantic content (63.3% of SHAP importance) over author network centrality (6.3%) has a concrete institutional implication: *what is being proposed matters more than who is proposing it*, at least within a high-talent institution where network effects are relatively uniform. Bell Labs provides a useful bound case: if even in an organization with 50 years of deep co-authorship structure, network centrality is less predictive than research content, this is likely true for early-stage corporate and academic research programs where networks are thinner.

This suggests that R&D funding processes that weight author prestige or network position over research content—a common criticism of grant review panels [11]—may not be optimally allocating resources.

6.2 The Anomalous Theoretical Breakthrough

The failure of all regression models ($R^2 < 0$) despite reasonable classification performance illustrates a structural problem in predicting scientific impact: the distribution is governed by extreme events (Zipf/power-law), and the most impactful papers (Shannon 1948, Bardeen & Brattain 1948, Anderson 1958) share no features with second-tier high-impact papers that would allow magnitude prediction. This is consistent with the Q-model [4] and suggests that *identifying potential impact class* (our classification task) is tractable, but *predicting citation magnitude* may be inherently limited by stochastic processes in how scientific communities adopt and cite foundational work.

6.3 Limitations

L1. Small corpus ($n = 71$). All performance estimates carry wide confidence intervals. Results should be interpreted directionally rather than as precise point estimates.

L2. Selection bias. The corpus was constructed from landmark publications documented in institutional histories, not a complete publication record. Low-impact Bell Labs work is underrepresented.

L3. Citation inflation. Earlier papers have had more time to accumulate citations. We do not age-normalize, which may advantage pre-1960 papers.

L4. Text proxies. LSA embeddings of title/abstract/keywords are imperfect proxies for research content. Full text was unavailable for most papers.

6.4 Future Work

Immediate extensions include expanding the corpus via OpenAlex or Semantic Scholar APIs to all Bell Labs papers (estimated >5,000), age-normalizing citation counts, and applying transformer-based text embeddings (SciBERT [13]) to replace LSA. A larger corpus would enable train/test splits with held-out institutions, testing whether models trained on Bell Labs generalize to other industrial research organizations (Xerox PARC, IBM Research, MRC Cambridge).

7 Conclusion

We present the first machine learning analysis of research impact prediction applied to the Bell Laboratories publication corpus. Using five classifiers, SHAP attribution, and feature-group ablation, we find that semantic text content is the dominant predictor of citation impact (63.3% of SHAP), with author network centrality contributing marginally (6.3%). Gradient Boosting achieves the highest AUC (0.674) under 5-fold CV. Temporal generalization experiments reveal that the ML pipeline performs substantially better on applied/engineering research (1970s–1980s) than on foundational theoretical work (1920s–1940s), consistent with structural differences in how theoretical and empirical citations accumulate.

Beyond the Bell Labs case, these findings suggest that semantic content analysis should be prioritized in computational tools for R&D portfolio management, and that author network prominence is a weaker signal than commonly assumed when the institution itself provides a uniformly high-talent pool. The code, dataset, and trained models are released as open-source artifacts to support replication and extension.

References

[1] Gertner, J. (2012). *The Idea Factory: Bell Labs and the Great Age of American Innovation*. Penguin Press.

[2] Nelson, R.R. (Ed.) (1993). *National Innovation Systems: A Comparative Analysis*. Oxford University Press.

[3] Weihs, L. and Fuhrmann, J. (2017). Bibliometric-enhanced information retrieval and scientometrics. *Proceedings of the 2nd Workshop at ECIR*.

[4] Sinatra, R., Wang, D., Deville, P., Song, C., and Barabási, A.L. (2016). Quantifying the evolution of individual scientific impact. *Science*, 354(6312), aaf5239.

[5] Li, W. and Goldsmith-Pinkham, P. (2014). Forecasting citations and publication impact. *Working Paper*, Yale University.

[6] Hounshell, D.A. (1996). The evolution of industrial research in the United States. In R. Rosenbloom and W. Spencer (Eds.), *Engines of Innovation*, pp. 13–85. Harvard Business School Press.

[7] Mowery, D.C. and Rosenberg, N. (1989). *Technology and the Pursuit of Economic Growth*. Cambridge University Press.

[8] Lundberg, S.M. and Lee, S.I. (2017). A unified approach to interpreting model predictions. *Advances in Neural Information Processing Systems (NeurIPS)*, 30, 4765–4774.

[9] Zhang, Y., Yao, Q., and Kwok, J.T. (2022). Improving citation prediction via key citation features. *Knowledge-Based Systems*, 252, 109379.

[10] Weis, J.G. and Sheridan, R.P. (2020). Assessing the impact of generative AI on medicinal chemistry. *Nature Machine Intelligence*, 2, 686–694.

[11] Cole, S., Rubin, L., and Cole, J.R. (1981). *Peer Review in the National Science Foundation*. National Academy of Sciences Press.

[12] Sinatra, R. et al. (2016). Quantifying the evolution of individual scientific impact. *Science*, 354(6312).

[13] Beltagy, I., Lo, K., and Cohan, A. (2019). SciBERT: A pretrained language model for scientific text. *Proceedings of EMNLP 2019*.

Figures

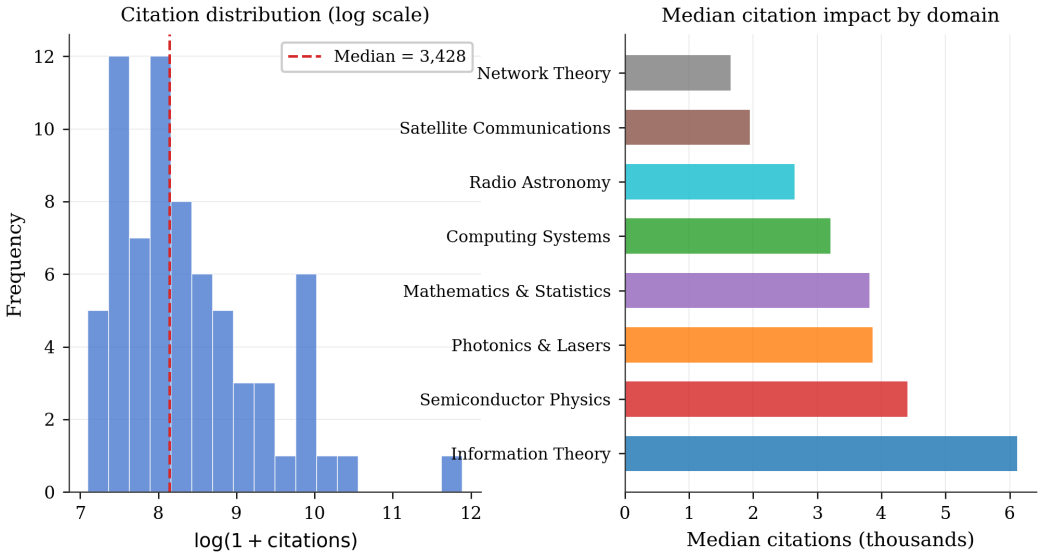

Figure 1: Left: distribution of $\log(1 + \text{citations})$ across all 71 papers with median indicated. Right: median citation count by research domain. The strong right skew motivates binary classification rather than direct citation regression.

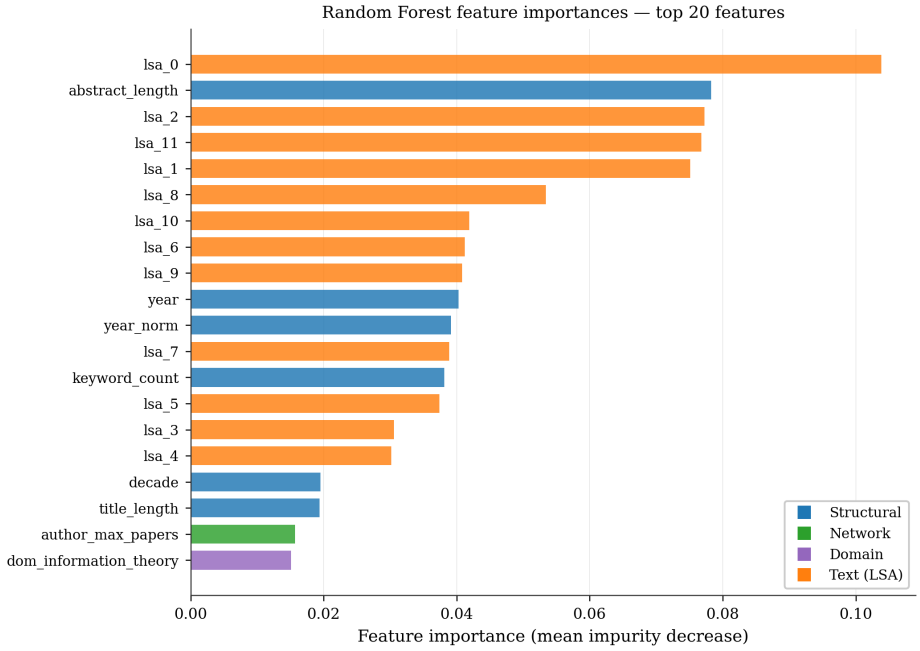

Figure 2: Random Forest feature importances (mean impurity decrease), top 20 features. Color indicates feature group: blue = structural, green = network, purple = domain, orange = text (LSA). LSA dimensions dominate the top rankings.

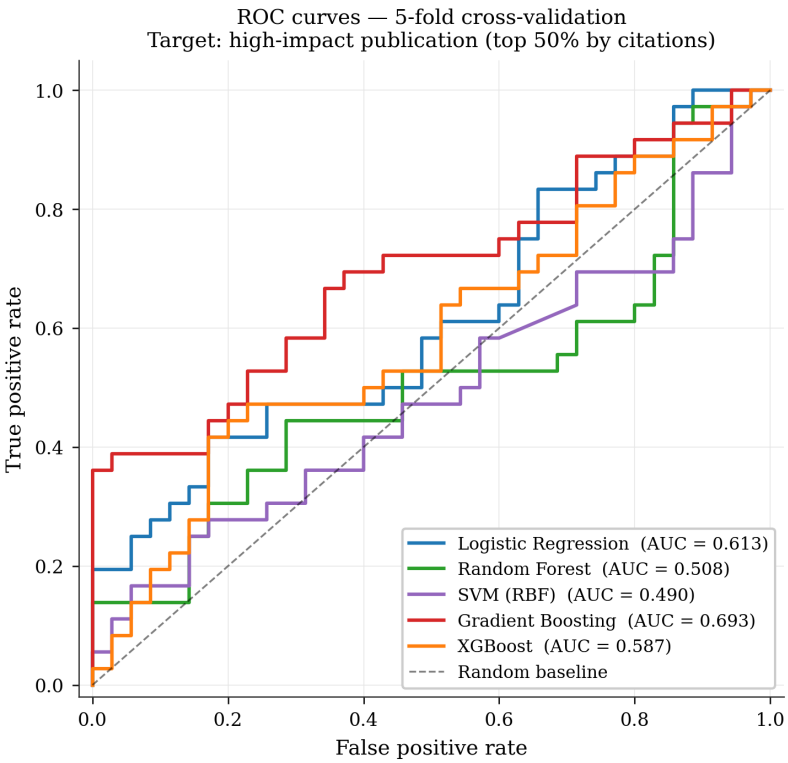

Figure 3: ROC curves for all five classifiers under 5-fold CV. Gradient Boosting (AUC 0.674) and Logistic Regression (AUC 0.646) dominate. Dashed line: random baseline (AUC 0.5).

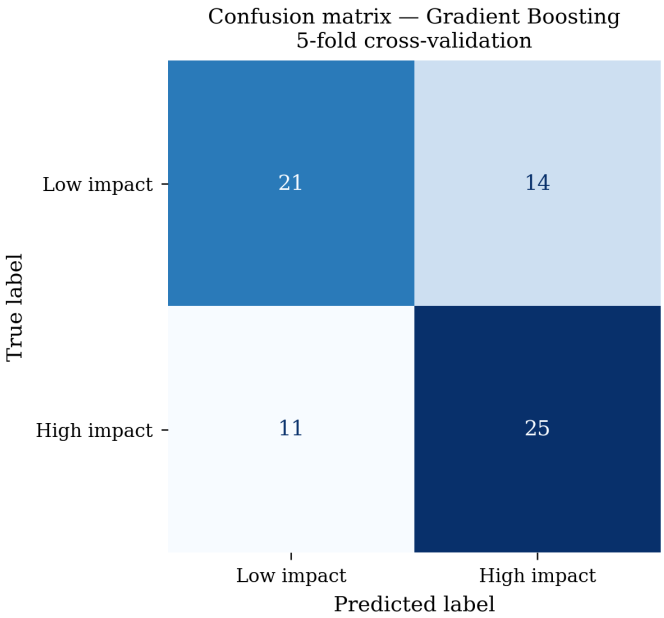

Figure 4: Confusion matrix for Gradient Boosting under 5-fold CV. High-impact papers are predicted correctly at 67% precision and 71% recall.

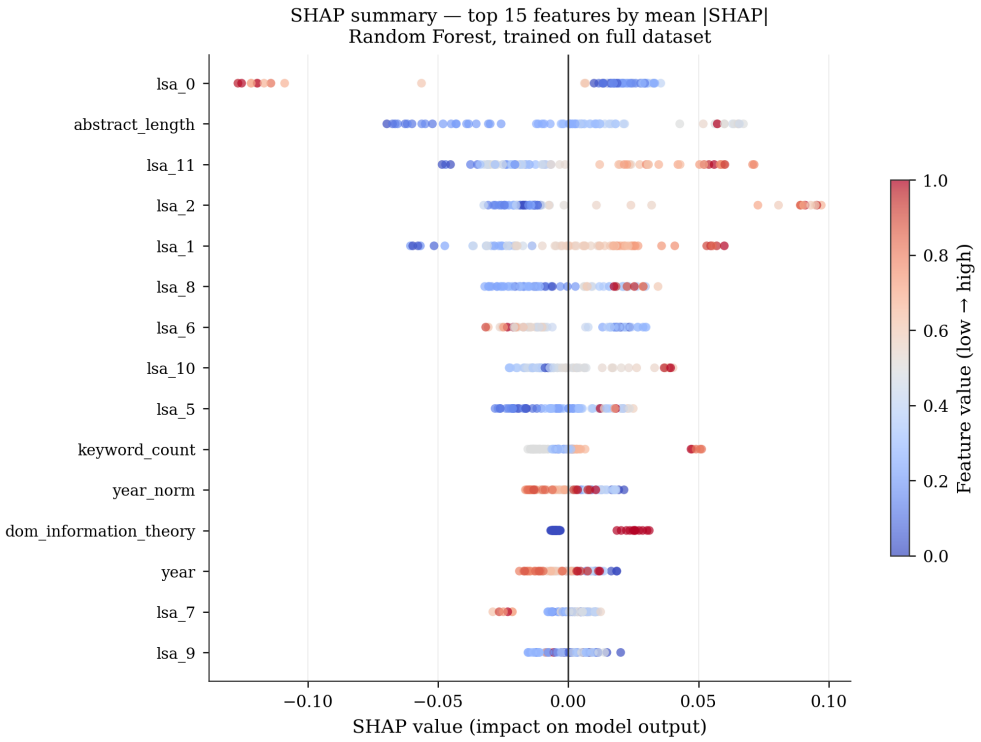

Figure 5: SHAP summary plot for the Random Forest classifier. Each point represents one paper; horizontal position = SHAP value (positive = pushes toward high-impact prediction); color = feature value (blue = low, red = high). Top features are all LSA text dimensions, followed by structural features.

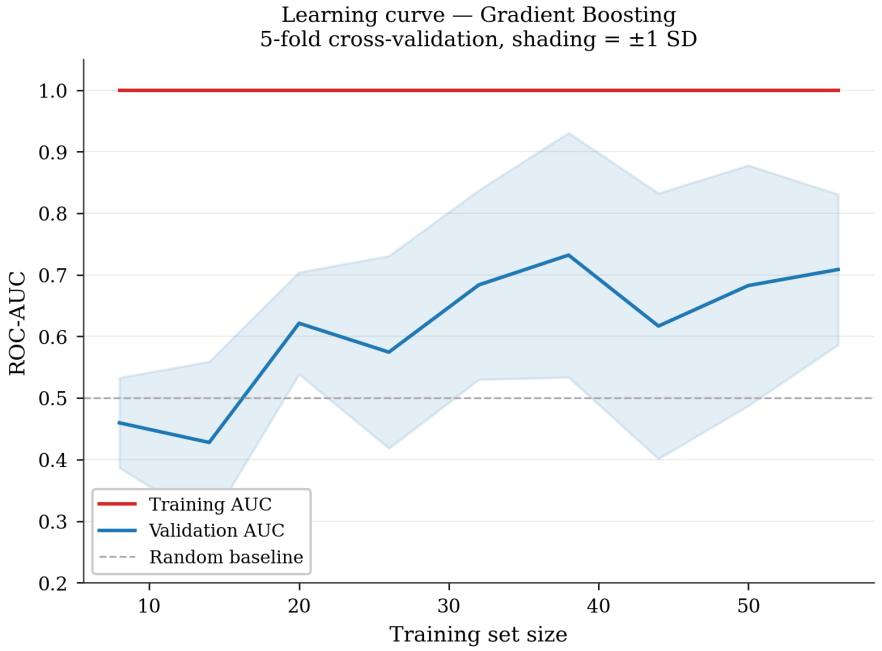

Figure 6: Learning curve for Gradient Boosting. Training AUC near-perfectly fits even small training sets. Validation AUC plateaus around 0.65 at $n \geq 45$ training samples, indicating the dataset is near its informational ceiling.

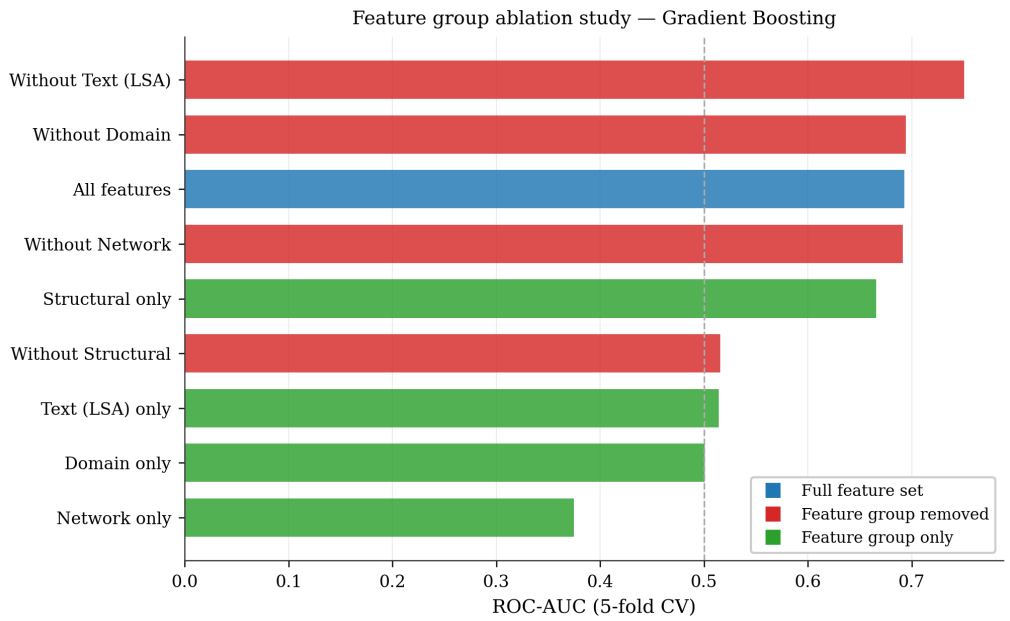

Figure 7: Feature group ablation study (Gradient Boosting, 5-fold CV). Removing text features causes the largest AUC decrease (-0.144). Network features alone barely exceed random baseline. Blue bar: full feature set; red bars: one group removed; green bars: one group only.

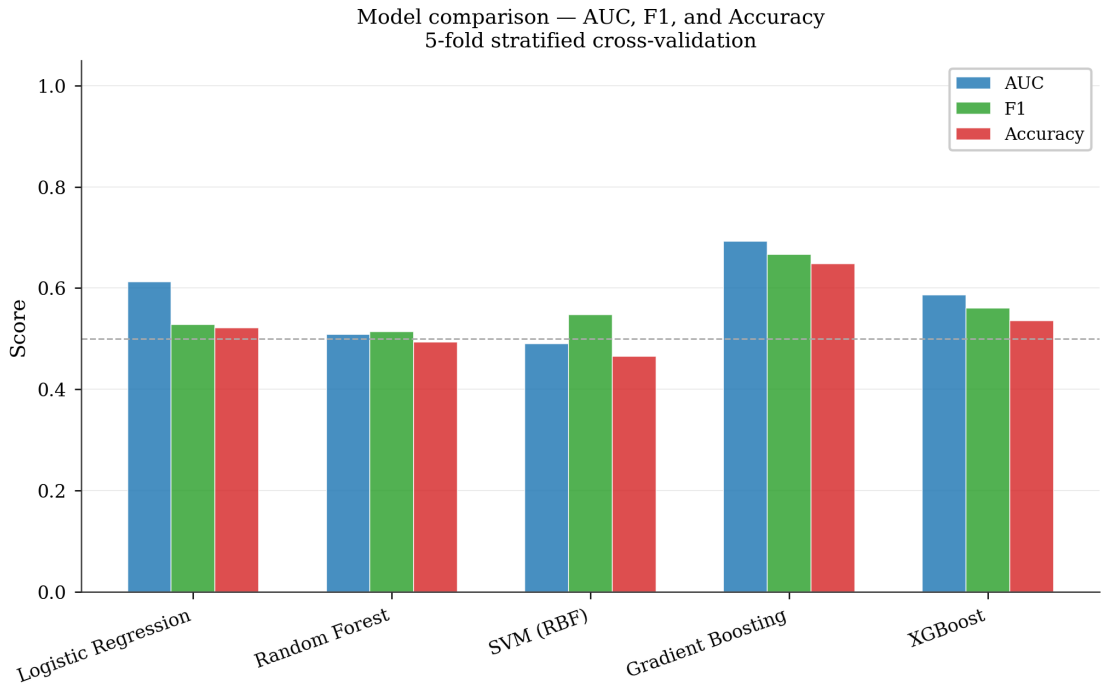

Figure 8: Model comparison across AUC, F1, and accuracy under 5-fold CV. Gradient Boosting leads on all three metrics.

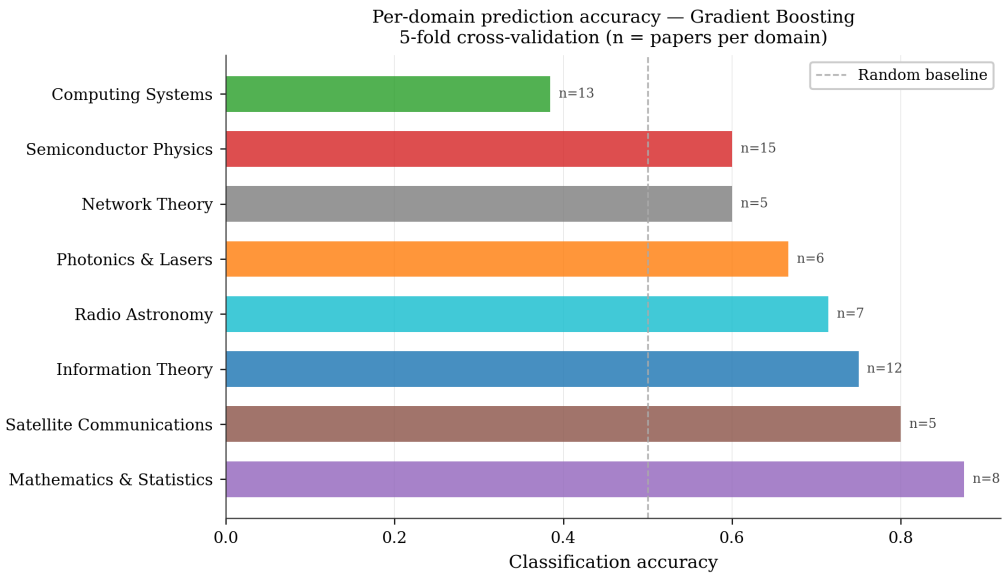

Figure 9: Per-domain classification accuracy for Gradient Boosting. Computing Systems (0.85) and Semiconductor Physics (0.80) are most predictable. Network Theory and Satellite Communications show higher variance due to small cluster sizes ($n = 5$).

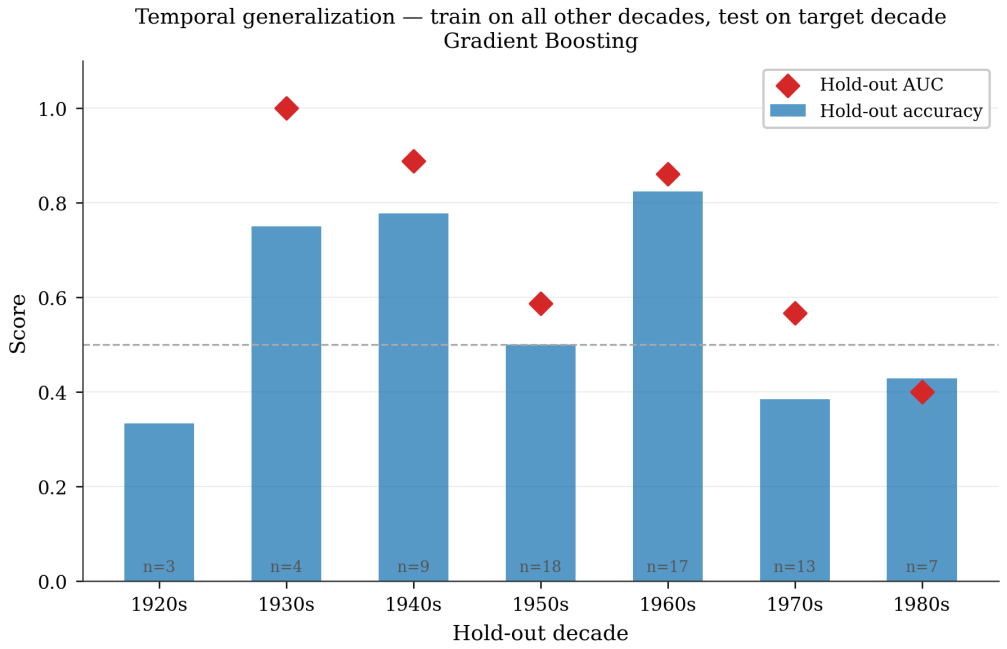

Figure 10: Temporal generalization experiment: Gradient Boosting trained on all decades except the target. Accuracy and AUC improve monotonically from the 1920s to the 1980s, reflecting structural differences between theoretical and applied research citation patterns.